

Hardware RSA Accelerator

Group 3: Ariel Anders, Timur Balbekov, Neil Forrester

May 13, 2013

Contents

1	Background	1
1.1	RSA Algorithm	1
1.1.1	Definition of Variables for the RSA Algorithm	1
1.1.2	Encryption	1
1.1.3	Decryption	2
1.2	RSA Algorithm for Hardware Implementation	2
1.2.1	Modular Exponentiation	2
1.2.2	Interleaved Modulus Multiplication	3
1.3	Implementation in C	3
1.3.1	Right to Left Binary Modular Exponentiator	5
1.3.2	Interleaved Modular Multiplier	5
1.4	Implemented Modules	8
1.4.1	Integer Representation Explorations	8
1.5	Challenges	9
1.5.1	Modular Exponentiation	9
1.5.2	Interleaved Modular Multiplication	9
1.5.3	Adder Implementation	9
1.6	Final Design Results: Space, Timing, and Throughput	10

Abstract

Our project is implementing the RSA cryptographic algorithm in Bluespec. The benefits of doing this in hardware are higher performance, reduced power usage and size, and cost. Having reusable IP that implements RSA would allow a device manufacturer to either reduce their power footprint, or skip inclusion of a processor in a device that otherwise would not need one.

More specifically, our objective is to implement the encryption and decryption protocols in RSA for 1024 bit messages on the XUPV5 Bluespec ...speed. (include some speed stuff here)

Chapter 1

Background

CITE WIKIPEDIA!! Assume the readers of your final report are not intimately familiar with your project domain. Provide enough high level background about the domain or algorithms involved for the reader to get an idea of the scope of the project and understand significant design choices described later in the document. Include the benefits of using an FPGA for this application.

RSA is an encryption protocol that involves a public and private key. As indignant of its name, the public key is known to everyone and is used to encrypt messages, which are called "ciphertext". The ciphertext can then be decrypted using the private key. Our implementation assumes the public and private keys are generated prior to encryption through with the FPGA.

1.1 RSA Algorithm

1.1.1 Definition of Variables for the RSA Algorithm

Public Key (n, e) A public key consists of the modulus n and encryption exponent e

Private Key (n, d) A private key consists of the same modulus n and the private decryption exponent d

Message (m) The plain text message converted into an integer m , such that $0 \leq m \leq n$

Ciphertext (c) The plain text message encrypted using the public key

1.1.2 Encryption

The ciphertext is generated by encrypting the plain text message using the public key:

$$c \equiv m^e \pmod{n} \quad (1.1)$$

1.1.3 Decryption

The plaintext message is generated by decrypting the ciphertext message using the private key:

$$m \equiv c^d \pmod{n} \quad (1.2)$$

1.2 RSA Algorithm for Hardware Implementation

Using a naive approach to implement this algorithm would clearly run quickly out of space. Instead, we have selected two algorithms that blablabla good space yada yada yada Copied algorithms over... still needs to be cleaned up

1.2.1 Modular Exponentiation

This module will employ the Right-to-left binary algorithm, which we believe is a good compromise between speed, memory usage, and complexity. The goal of the algorithm is to calculate $b^e \bmod m$ for very large values of b , e , and m . If the bits of e are $e_1, e_2 \dots e_n$:

$$e = \sum_{i=0}^n e_i 2^i \quad (1.3)$$

then:

$$b^e = \prod_{i=0}^n e_i b^{(2^i)} \quad (1.4)$$

and since:

$$a * b \bmod m = (a \bmod m) * (b \bmod m) \bmod m \quad (1.5)$$

then every intermediate result can be taken modulo m to keep the size of intermediate results manageable. Therefore, the following algorithm will compute $b^e \bmod m$ in a reasonable amount of time and memory:

b , e , and m are the inputs to the algorithm.

$c \leftarrow 1$

while $e > 0$ **do**

if $e \bmod 2 = 1$ **then**

$c \leftarrow c * b \bmod m$

end if

$b \leftarrow b * b \bmod m$

$e \leftarrow \lfloor e/2 \rfloor$

end while

c is the result of the algorithm.

This very naturally suggests a circular pipeline in hardware. If parallelism is desired, then multiple circular pipelines may be put in parallel, with some

logic at the front and back to manage handing out jobs to different circular pipelines, and collecting the results. The only remaining problem is performing multiplication, modulo, and bit shifting which is implemented through the interleaved modular multiplication algorithm which requires bit shifts, additions, subtractions, and bitwise comparisons and does not take up excessive area.

1.2.2 Interleaved Modulus Multiplication

N is the size of the numbers, in bits. For example, $N = 1024$. Also, x_i is the i th bit of x .

x , y , and m are the inputs to the algorithm.

$p \leftarrow 0$

$i \leftarrow N - 1$

while $i \geq 0$ **do**

$p \leftarrow p * 2$

if $x_i = 1$ **then**

$p \leftarrow p + y$

end if

if $p \geq m$ **then**

$p \leftarrow p - m$

end if

if $p \geq m$ **then**

$p \leftarrow p - m$

end if

$i \leftarrow i - 1$

end while

p is the result of the algorithm.

1.3 Implementation in C

We have a working implementation of all our algorithms in C, that we wrote from scratch and verified its correctness with libgcrypt. The C implementation, of course, is unable to operate directly on 1024 bit integers, so we store them as arrays of 16 bit unsigned integers. As a result, performing bit shifts, additions, and comparisons takes somewhat more code than it would take to perform the corresponding operations in Bluespec.

High-Level Design and Test Plan

Our RSA module has three inputs: data, exponent, and modulus; this is because the only difference between implementing encryption versus decryption is the inputs to our RSA unit. For example, to encrypt a plain text message the input to our RSA unit is the message, the public key and its modulus (in this case the "data" is the plain text message and the "exponent" is the public key). The output of this unit is the encrypted ciphertext. Alternatively, to decrypt the cipher text, the input to our RSA unit is the ciphertext, the private key and its modulus. Figure bla represents the high level design of our RSA unit. In our implementation, all bit lengths are 1024 bits long and the input and output to the FPGA is implemented using the SceMi test bench.

To verify the functionality of the RSA module, we compared the results of the encryption and decryption blocks to the results of a software implementation. The two private keys for encryption and decryption modules are passed by SceMi into the hardware. A SceMi testbench pushes a message to the encryption block, along with an enable signal, message, and public key of the software test-bench. The module generates an encrypted message, and the software test-bench uses its private key to decrypt and verify the correctness of the encrypted message.

For decryption, the process is reversed: the test-bench passes in an encrypted message instead of plain-text, and the decryption module uses the private key of the software test-bench to decrypt the message. Then, the test-bench verifies the plain-text for correctness.

Microarchitectural Description

Our project is divided into two important modules: `ModExpt.bsv` and `ModMultIllvd.bsv`. `ModExpt.bsv` performs modular exponentiation, while `ModMultIllvd.bsv` performs modular multiplication using the interleaved modular multiplication algorithm described above. The modular exponentiator instantiates two modular multipliers. The high level diagram in Figure 1.1 depicts the interface between the modular multipliers and the modular exponentiator.

1.3.1 Right to Left Binary Modular Exponentiator

The modular exponentiator is a circular pipeline (depicted in Figure 1.2). On each cycle of the pipeline it supplies inputs to the two multipliers. When the multipliers complete, it stores the results back into the registers. However one result is discarded if the low bit of e is 0. In fact, our actual implementation will probably not invoke the multiplier if its result will be discarded anyway. However, this is simply an optimization, and doesn't hugely affect the overall plan. On every iteration, the value of e is right-shifted by one bit. When e is zero, the loop terminates.

1.3.2 Interleaved Modular Multiplier

The interleaved modular has the advantage of not requiring long multiplies, and works with only left shifts, addition, subtraction, and comparison. Unfortunately, a step of the algorithm requires comparing the entire length of the data in the worst case. Additionally, there are 3 possible add/subtract steps at every step of the algorithm. Therefore, the propagation delay of each step of the algorithm is prohibitive without pipelining. The naive, un-pipelined approach did meet timing because of the long propagation delay through the adders.

An overview of the module is pictured in Figure 1.3.

RSA

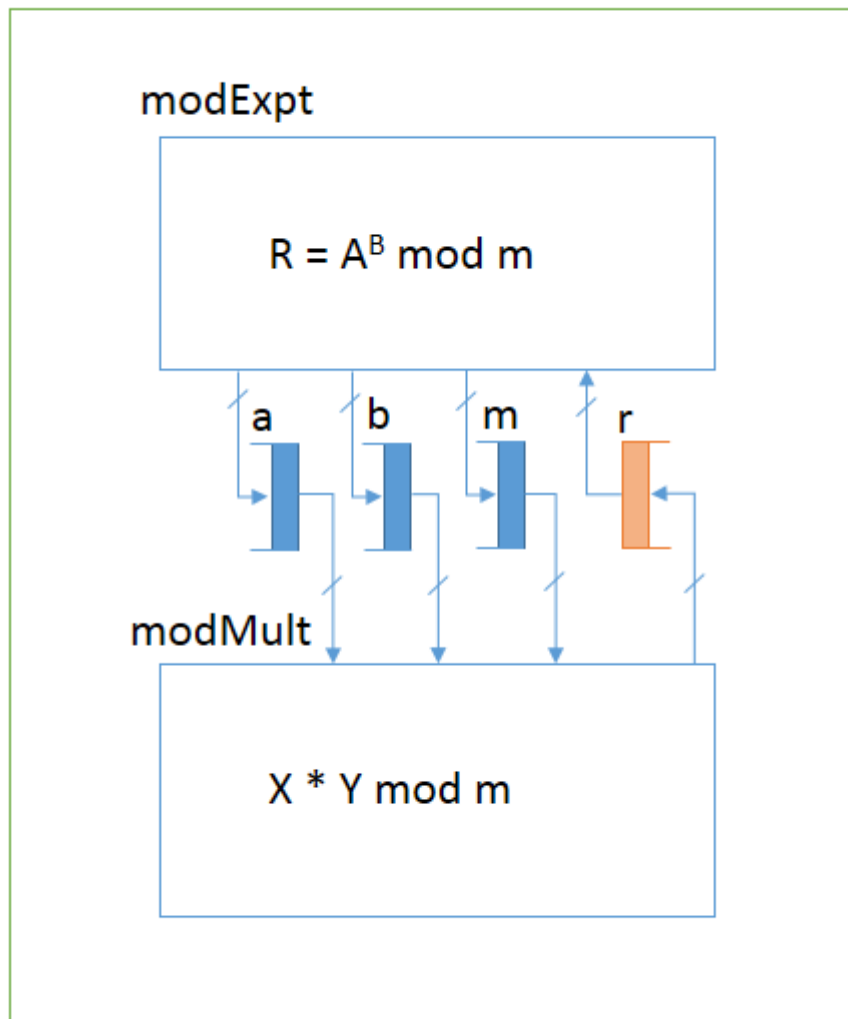


Figure 1.1: High level overview. Note that only one of two multipliers is depicted.

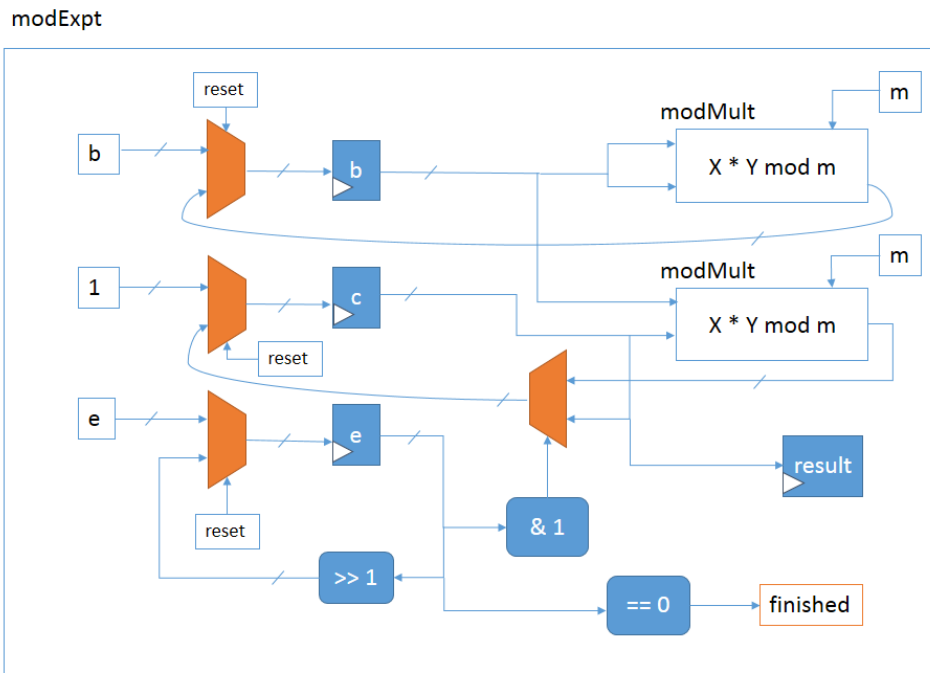


Figure 1.2: Modular exponentiation

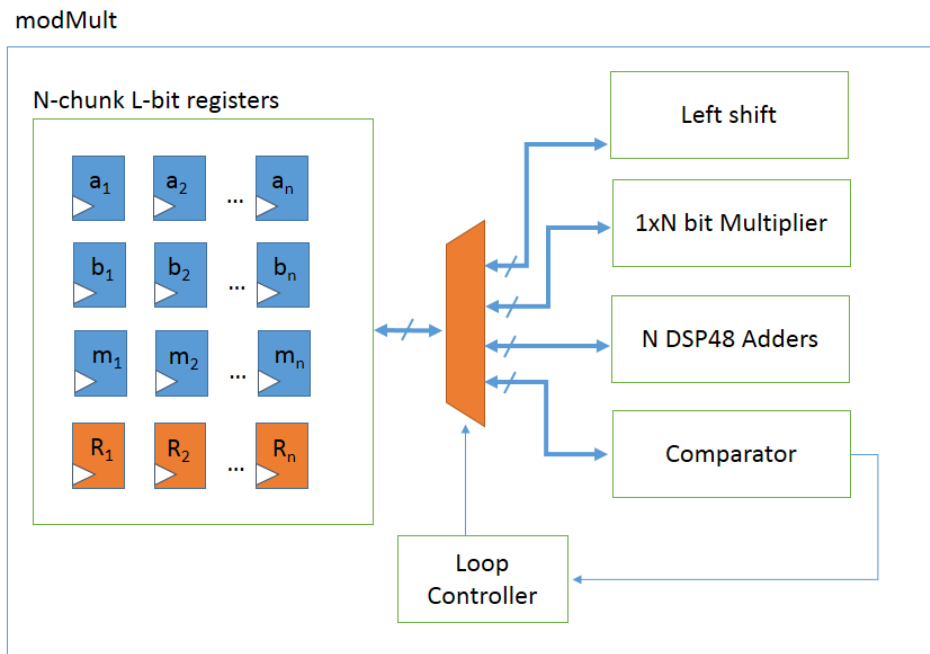


Figure 1.3: Interleaved modular multiplication

1.4 Implemented Modules

The modules we implemented are: `SceMiLayer.bsv`, `RSAPipeline.bsv`, `RSA.bsv`, `ModMultIlvd.bsv`, `ModExpt.bsv`, `PipelineAdder.bsv`, `CLAdder.bsv`

SceMiLayer.bsv We developed two SceMiLayer modules to test different types of designs: one for importing vmh files, and another for importing libgcrypt for simulating our c code.

RSAPipelinetypes.bsv This is the general header file where we define all constant values. This will make the overall product modular and easy for others to change core elements of design such as number of bits per chunk.

RSA.bsv This is a dedicated alternative driver for the RSA module that performs cosimulation with libgcrypt

ModMultIlvd.bsv The interleaved modulus multiplier based on the Montgomery Algorithm. This function computes $a * b \bmod m$.

ModExpt.bsv The modulus exponentiator implements the algorithm described above. It creates two modulus multiplier to computer $b^e \bmod m$.

PipelineAdder, CLAdder.bsv These modules implement a folded and carry look-ahead adders respectively.

1.4.1 Integer Representation Explorations

1. The first interface was a simple `Int#(1024)` representation. We created a simple adder in order to synthesize our design.
2. The second type of interface is most similar to our C implementation where integers are stored as 64 - 16 bit chunks.
3. The third type of interface uses BRAM to store chunks of the integer throughout the implementation. (This is currently incomplete)

Difficulties Encountered

We created a simple adder in order to synthesize our design. Since we had doubts about the success of this representation this was a vital step before continuing our design. Our concerns were well-founded for the simple `Int#(1024)` representation: the simple addition of two `Int#(1024)` were unable to synthesize.

Implementation Evaluation

1.5 Challenges

Challenges: meeting timing and space constraints on the fpga. We made significant changes to our original microarchitecture to successfully meet our initial project goals. - we fixed timing by modifying the modular multiplication module and space modifying our modular exponentiation architecture.

1.5.1 Modular Exponentiation

blbabla write about space constraints and using one multiplier only include image of modexpt diagram

1.5.2 Interleaved Modular Multiplication

write about using 1 adder and 2 subtraction units. To increase the cycle time we do 2 subtractions: $p-m$ and $p-2m$ and select the non negative result at the end. include diagram

1.5.3 Adder Implementation

Adder was on the critical path

Clader

Multi-cycle Adder

existing ip blocks used: bram and +

Synchrhon

existing ip blocks used

1.6 Final Design Results: Space, Timing, and Throughput

What was the device utilization of your design, the clock frequency, and the high level throughput (frames per second, bits decoded per second, etc...)? Did you meet your initial project goals? Why or why not?

Design Exploration

Now that you have a working system, what trade-offs would you be interested in exploring? If you didn't meet your performance targets, why? What are the limiting components in your design? A detailed analysis of the cycle latency and throughput of each component in your design may be appropriate.